

Lecture 4: Transformers and pretraining-finetuning

CS6493 Natural Language Processing
Instructor: Linqi Song and Chen MA



Outline

- 1. Attention
- 2. Transformer
- 3. BERT
- 4. GPT

Problems with contextualized word embeddings

- RNNs/LSTMs have **long-term dependency** problems, where words/tokens are processed sequentially.
 - Information loss
 - Hard to compute in parallel
- Bi-directional RNNs/LSTMs **feature fusion** and **representation** ability is weak (compared with transformers).

From ELMo to BERT

ELMo



Embeddings from Language Models

BERT



Bidirectional Encoder Representations from Transformers

Attention – the main technique behind transformers

- Attention is all you need, Google, 2017.
- Why attention?
 - To reduce the computational **complexity**.
 - To **parallelize** the computation.
 - Self-attention connects all positions in a sequence with a constant number of operations to solve the **long-term dependency** problem.
 - Self-attention could yield more **interpretable** models





What is attention?

- What is the attention mechanism
 - An analogy to human's brain, to pay more attention to more important information.
 - Map a **query** and a set of **key-value pairs** to an **output**, where the query, keys, values and output are all vectors.
 - The output is a **weighted sum of the values**, where the weight assigned to each value is computed by a compatibility function of the query with the corresponding key.




Attention in cognitive psychology and graphical design

- Cognitive psychology and graphical design
 - Which part of the screen attracts the most human attention: ad placement, news headlines, etc.
 - Visual hierarchy




Attention in NLP: aligning while translating (1)

- Machine translation: translate source sentence in a language to a target sentence in another language.




- Alignment in machine translation: relationship between source words and target words.




Attention in NLP: aligning while translating (2)

- Recall the RNN-based encoder-decoder example for neural machine translation



Attention in NLP: aligning while translating (3)

- Attention intuition: each time the proposed model generates a word in a translation, it searches for a set of positions in a source sentence where the most relevant information is concentrated (Bahdanau et al. 2015).



- Do not encode all inputs into one vector: more information.
- Allow adaptive selection to which should the model attend.

Decoder RNN

Attention: formal description

- Given a **query** vector $\mathbf{q}$, and a set of **key-value** pairs (all vectors) $\{(\mathbf{k}_i, \mathbf{v}_i)\}_{i=1}^L$, we first calculate the **similarity/attention score** between the query and each key:

$$s_i = \text{similarity}(\mathbf{q}, \mathbf{k}_i).$$

- Normalize the similarity score to be between 0 and 1, and they sum up to 1. These are called **attention distribution**. One way is to use the softmax operation.


$$a_i = \text{softmax}(s_i) = \frac{\exp(s_i)}{\sum_{i=1}^L \exp(s_i)}.$$

- Compute the **attention/context vector** $\mathbf{z}$ as a weighted sum of values.

$$\mathbf{z} = \sum_{i=1}^L a_i \mathbf{v}_i.$$

- Keys and values are not necessarily the same, and they could be different as well, such as in machine translations.

Attention: example



Attention: similarity calculation

- Similarity/attention score calculation between the query and each key:

$$s_i = \text{similarity}(\mathbf{q}, \mathbf{k}_i).$$

- In general, every similarity measures can be used here, such as the cosine coefficient and Pearson correlation coefficient.

- Commonly used ones for neural networks:

- Additive attention (Bahdanau et. al 2014):

$$s_i = w_3 \tanh(\mathbf{w}_1^T \mathbf{q} + \mathbf{w}_2^T \mathbf{k}_i).$$

- Multiplicative attention.

$$s_i = \mathbf{q}^T \mathbf{W} \mathbf{k}_i, \text{ where } \mathbf{W} \in \mathbf{R}^{d_q \times d_k}, \text{ with } d_q, d_k \text{ being the dimensions of } \mathbf{q} \text{ and } \mathbf{k}_i.$$

- Dot-product attention.

$$s_i = \mathbf{q}^T \mathbf{k}_i.$$

- Scaled dot-product attention.

$$s_i = \frac{\mathbf{q}^T \mathbf{k}_i}{\sqrt{d_k}}.$$

Scaled dot-product attention in matrix form

Q : a matrix formed by packing a set of queries


K : a matrix formed by packing a set of keys

V : a matrix formed by packing a set of values

d_k : the dimension of keys and queries



$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Scaled Dot-Product Attention




Attention learns (nearly) alignment

- An example from the paper 'Neural Machine Translation by Jointly Learning to Align and Translate'




Attention is all you need: self-attention (1)

- For an input sequence of words, play attention mechanisms between every word and others (including itself).
- Features
 - Constant path length between any two positions
 - Easy to parallelize per layer



Attention is all you need: self-attention (2)

- How to construct queries, keys, and values?
 - Linear transformation from original input embeddings x_i : $q_i = x_i W^Q$, $k_i = x_i W^K$, $v_i = x_i W^V$



Attention is all you need: self-attention (3)

- Perform scaled dot-product attention between a word i and all words
- Obtain L self-attention vectors, one for each word (token).

Input

Embedding

Queries

Keys

Values

Score

Divide by 8 ($\sqrt{d_k}$)

Softmax

Softmax
X
Value

Sum

Thinking

x_1

q_1

k_1

v_1

$q_1 \cdot k_1 = 112$

14

0.88

v_1

z_1

Machines

x_2

q_2

k_2

v_2

$q_1 \cdot k_2 = 96$

12

0.12

v_2

z_2


Attention is all you need: self-attention (4)

- We change these calculations into the matrix form.

$$\mathbf{X} \times \mathbf{W}^Q = \mathbf{Q}$$



$$\mathbf{X} \times \mathbf{W}^K = \mathbf{K}$$


$$\mathbf{X} \times \mathbf{W}^V = \mathbf{V}$$


$$\text{softmax}\left(\frac{\mathbf{Q} \times \mathbf{K}^T}{\sqrt{d_k}}\right) \mathbf{V}$$

$$= \mathbf{Z}$$

Attention is all you need: multi-head self-attention (1)

- What if we choose different forms of W^Q , W^K , and W^V ?
 - Each group of (W^Q, W^K, W^V) is called a head.
 - Increase the representation performance: different representation subspaces, model's ability to focus on different positions.



Attention is all you need: multi-head self-attention (2)

- How to combine multi-headed output together?
 - Concatenate them and use another linear transformation.


1) Concatenate all the attention heads



2) Multiply with a weight matrix W^O that was trained jointly with the model

X

3) The result would be the Z matrix that captures information from all the attention heads. We can send this forward to the FFNN



Attention is all you need: multi-head self-attention (3)

- In a nutshell



1) This is our input sentence*

Thinking
Machines


2) We embed each word*




3) Split into 8 heads.
We multiply X or R with weight matrices




...



4) Calculate attention using the resulting $Q/K/V$ matrices



...



5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer




...




W^O



Z



* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one




Attention is all you need: multi-head self-attention (4)

- Graphical view and mathematical presentations
 - Linearly project the queries, keys and values h times with different, learned linear projections to d_k, d_k, d_v dimensions respectively.
 - On each of these projected versions, perform the attention function in parallel.
 - The resulting d_v - dimensional output values are concatenated and once again projected.
- Parameters: hidden size d_{model} , self-attention heads h

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W^O$$



where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$

Where the projections are parameter matrices $W_i^Q \in \mathbb{R}^{d_{\text{model}} \times d_k}$, $W_i^K \in \mathbb{R}^{d_{\text{model}} \times d_k}$, $W_i^V \in \mathbb{R}^{d_{\text{model}} \times d_v}$ and $W^O \in \mathbb{R}^{hd_v \times d_{\text{model}}}$.





Multi-head attention results illustration

- One head attention for 'it' -> more on 'the animal';
- Multi-head attention for 'it' -> 'street', 'tire', 'didn't', 'cross', 'because'.




Transformer

- Proposed by Google: Attention is All You Need (Vaswani et al., 2017)
- Main technique: [multi-head self attention mechanism](#).
- The transformer is a novel architecture that aims to solve [sequence-to-sequence](#) tasks while handling long-range dependencies with ease.




Transformer - formal description

- Encoder: maps an input sequence of symbol representations $(x_1, x_2, \dots, x_n)$ to a sequence of continuous representations $z = (z_1, z_2, \dots, z_n)$
- Decoder: given z , generates an output sequence $(y_1, y_2, \dots, y_m)$, one element at a time.
- At each step the model is **auto-regressive**, consuming the previously generated symbols as additional input.
- Stacked **self-attention** and **point-wise, fully connected layers** for both the encoder and decoder.



Transformer: attention in encoder

Encoder




- In a self-attention layer all of the **keys, values and queries** come from the **output of the previous layer or the raw input embedding** (with hidden dimension d_{model}) in the encoder.
- Each position in the encoder can attend to all positions in the previous layer of the encoder.

Transformer: attention in decoder

Decoder

- Similarly, each position attends to all positions in the decoder up to and including that position.
- **Encoder-decoder attention**: **keys and values** come from the top **encoder** output; **queries** come from the output of the masked multi-head attention (with hidden dimension d_{model}) in the **decoder**.
- **Masked self-attention**
 - **Keys, values and queries** come from the **output of the previous layer or the raw output embedding** (with hidden dimension d_{model}) in the decoder.
 - **Prevent leftward information flow** to preserve the auto-regressive property. Decoder self-attention is only allowed to attend to earlier positions in the output sequence, implemented by **masking out** (setting to $-\infty$) future positions (tokens).




Transformer: other details

- Each of the layers in the encoder and decoder contains a fully connected feed-forward network (FFN), which is applied to each position separately and identically.

$$FFN(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

- Positional encodings**: input positional information of the sequence.
- Residual connection: learning the changing part
- Learned embeddings** are used to convert the input tokens and output tokens to vectors of dimension d_{model} . The same weight matrix are shared in the two embedding layers.



Positional encoding in transformer (1)

- Transformer does not contain recurrence or convolution, it does not know the order of input tokens.
- We have to let the model know the positions of the tokens explicitly
- Idea: input representation of a token is the sum of two embeddings: token and positional

Positional encoding in transformer (2)


- In order to add position information (order of the sequence)

$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$

- pos is position (e.g., the pos -th token in a sentence) and i is the i -th element in a vector to represent each position.
- Each dimension of the positional encoding corresponds to a sinusoid.
- For any fixed offset k , PE_{pos+k} can be represented as a linear transformation of PE_{pos} . This would allow the model to easily learn to attend by relative positions.

Positional encoding in transformer (3)



BERT: Bidirectional Encoder Representations from Transformers

- BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding, Google, 2018
- Pretrained deep **bidirectional** representation models from unlabeled text .
- Jointly conditioning on both left and right context in all layers.
- Can be finetuned with just one additional output layer to create models for a wide range of tasks.
- Powerful in a variety of tasks.

Task	Score	Best with BERT (from Delvin et al.)	Absolute improvement
NLU	GLUE score	80.5%	7.7%
NLI	MultiNLI accuracy	86.7%	4.6%
question-answer	SQuAD v1.1 Test F1	93.2	1.5
question-answer	SQuAD v2.0 Test F1	83.1	5.1


BERT architecture (1)

- BERT architecture consist of multi-layer bidirectional transformer encoders
- BERT model provided in the paper came in two sizes

BERT_BASE	BERT_LARGE
Layers = 12	Layers = 24
Hidden size = 768	Hidden size = 1024
Self-attention heads = 12	Self-attention heads = 16
Total parameters = 110M	Total parameters = 340M


BERT architecture (2)

- A graphical view



BERT: pretraining + finetuning

- Apart from output layers, the same **architectures** are used in both pre-training and fine-tuning.
- The **same pre-trained model parameters** are used to initialize models for **different down-stream tasks**.
- During **fine-tuning**, **all parameters** are fine-tuned.




What contributes to BERT's success?

- Innovation in model pre-training: 2 self-supervised tasks for pre-training instead of traditional unidirectional language models
 - Masked Language Model (MLM)
 - Next Sentence Prediction (NSP)
 - The training loss is the sum of the mean MLM likelihood and the mean NSP likelihood
- Compatibility to various tasks
 - Just fine-tune BERT Model for specific tasks to achieve state-of-the-art performance
 - BERT advances the state-of-the-art for eleven NLP tasks

BERT: input/output representation

- For a given token, its input representation is constructed by summing the corresponding token, segment, and position embeddings.
- [CLS] is a special symbol added in front of every input example.
- [SEP] is a special separator token (e.g. separating questions/answers).
- Unambiguously represent both a single sentence and a pair of sentences in one token sequence.



Pre-training BERT


- In order to pre-train BERT transformers bidirectionally, two self-supervised tasks are proposed, namely **Masked LM** and **Next Sentence Prediction (NSP)**.
- One of the reasons of BERT's success is the amount of data that it got trained on.
- The two corpora that were used to pre-train the language models are:
 - The BooksCorpus (800M words)
 - English Wikipedia (2,500M words)

Masked language modeling (MLM)

- Motivation: Enable bidirectional pre-training
- Main idea: Mask some percentage of the input tokens at random, and then predict those masked tokens.
- Details:
 - The training data generator chooses 15% of the token positions at random for prediction.
 - If the i -th token is chosen, we replace the i -th token with
 - (1) the [MASK] token 80% of the time,
 - (2) a random token 10% of the time,
 - (3) the unchanged i -th token 10% of the time.
 - Predict the original token via cross entropy and the final hidden vector for the i -th input token
- The masking scheme here is to reconcile the mismatch between pre-training and fine-tuning, since the [MASK] token does not appear during fine-tuning.


An illustration of MLM

Use the output of the masked word's position to predict the masked word




Next sentence prediction (NSP)

- Motivation: understanding the relationship between two sentences (inspired by Question Answering and Natural Language Inference).
- Main idea: given a preceeding sentence A and a sentence B, judge if B is the next sentence of A.
- Details:
 - choose the sentences A and B for each pretraining example
 - 50% of the time B is the actual next sentence that follows A (labeled as [IsNext])
 - 50% of the time it is a random sentence from the corpus (labeled as [NotNext])
- In the architecture figure, C is used for NSP prediction.



An illustration of NSP



NSP example

Input = [CLS] the man went to [MASK] store [SEP]
 he bought a gallon [MASK] milk [SEP]

Label = **IsNext**

Input = [CLS] the man [mask] to the store [SEP]
 penguin [MASK] are flight ##less birds [SEP]

Label = **NotNext**



BERT: fine-tuning (1)

- Common approach: independently encode text pairs before applying attention.
- BERT's approach: concatenate a text pair, and then encode with self-attention in order to include bidirectional cross attention between two sentences.
- For each task, simply plug in the task-specific inputs and outputs into BERT and fine-tune all the parameters end-to-end.
- Compared to pre-training, fine-tuning is relatively inexpensive.

BERT: fine-tuning (2)

- For **classification** tasks (e.g. sentiment analysis) we add a classification FFN for the [CLS] token input representation on top of the final output
- For **Question Answering** alike tasks BERT train two extra vectors that are responsible for marking the beginning and the end of the answer

BERT fine-tuning for classification




BERT fine-tuning for different tasks


- Context vector $\mathcal{C}$: Take the final hidden state corresponding to the first token in the input: [CLS].
- Transform to a probability distribution of the class labels:

$$P = \text{softmax}(CW^T)$$


- Batch size: 16, 32
- Learning rate (Adam): 5e-5, 3e-5, 2e-5
- Number of epochs: 3, 4




(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG



(b) Single Sentence Classification Tasks:
SST-2, CoLA



(c) Question Answering Tasks:
SQuAD v1.1



(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

Experiments on BERT

BERT fine-tuning results on 11 NLP tasks

Task	Score	Best with BERT (from Delvin et al.)	Absolute improvement
NLU	GLUE score	80.5%	7.7%
NLI	MultiNLI accuracy	86.7%	4.6%
question-answer	SQuAD v1.1 Test F1	93.2	1.5
question-answer	SQuAD v2.0 Test F1	83.1	5.1

Generative Pre-training (GPT)

- Motivation: semi-supervised learning
 - Stage 1: Unsupervised learning on word-level or phrase-level
 - E.g. word embeddings
 - Stage 2: Supervised training using these word-level features
- In the method above:
 - Stage 1 is **less dependent** on the task - related more to NLU
 - Stage 2 models may vary a lot according to different tasks
- Are there any techniques that improve the performance of Stage 1 models that would work for most of the tasks?
 - That's the problem we hope to solve by introducing pre-training.
 - In other words, pre-training improves **language understanding**.


GPT framework (1)

- Multi-layer transformer decoder
 - first layer: $h_0 = UW_e + W_p$
 - the l -th layer: $h_l = \text{transformer_block}(h_{l-1}), \forall l \in [1, n]$
 - Self-supervised pre-training, similar to embeddings such as word2Vec
 - Given tokens U , maximize the contextual conditional probability

$$L_1(U) = \sum_i \log P(u_i | u_{i-k}, \dots, u_{i-1}; \Theta)$$
$$P(u) = \text{softmax}(h_n W_e^T)$$

where k is the window size; $h_n W_e^T$ is the score for each word.

- Unidirectional!



GPT framework (2)


- Supervised fine-tuning
 - keep the pre-trained transformers
 - replace the final linear layer W_e with W_y ,
 - Given data inputs X and labels y , maximize

$$P(y|x^1, \dots, x^m) = \text{softmax}(h_l^m W_y).$$

$$L_2(\mathcal{C}) = \sum_{(x,y)} \log P(y|x^1, \dots, x^m).$$

$$L_3(\mathcal{C}) = L_2(\mathcal{C}) + \lambda * L_1(\mathcal{C})$$

auxiliary training objective



Task-specific adaptations (1)

How to adapt a single architecture to multiple input formats?

Classification (e.g. sentiment analysis)

- Given a piece of text, is it positive or negative?
- Answers: "Yes", "No"
- Answers: "Very positive", "Positive", "Neutral", "Negative", "Very negative"

Entailment

- Given a premise p and a hypothesis h , does p imply h ?
- Answers: "entailment", "contradiction", or "neutral"


Similarity

- Are two sentences semantically equivalent?
- Answers: "Yes", "No"

Multiple Choice (e.g. Story Cloze)

- Given a short story and two sentences, which is the sentence that ends the story?
- Given a passage and a question, and some multiple-choice answers, which is the answer?
- Answers: $A_1, A_2, \dots, A_N$

Task-specific adptions (2)



Pre-trained model comparison: ELMo, GPT, and BERT

	Architecture	Pre-training	Downstream tasks
ELMo	Bi-directional LSTM language model	Unsupervised corpus. Learn both words and linguistic context features that support downstream tasks.	Feature-based tasks and task-specific models.
GPT	Uni-directional transformer decoder	Unsupervised corpus. Each specific task requires discriminative fine-tuning .	Model-based tasks and task-agnostic models.
BERT	Bi-directional transformer encoder	Unsupervised corpus. 2 unsupervised tasks: MLM and NSP. Each specific task requires discriminative fine-tuning .	Model-based tasks and task-agnostic models.

References

Attention and transformer:

- Attention is All You Need (Vaswani et al., 2017) <https://arxiv.org/pdf/1706.03762.pdf>
- https://lena-voita.github.io/nlp_course/seq2seq_and_attention.html
- Neural Machine Translation by Jointly Learning to Align and Translate (Bahdanau, 2015) <https://arxiv.org/pdf/1409.0473.pdf>
- Transformer: <http://jalammar.github.io/illustrated-transformer/>

GPT:

- Improving Language Understanding by Generative Pre-Training (Radford et al., 2018) <https://www.cs.ubc.ca/~amuham01/LING530/papers/radford2018improving.pdf>

BERT:

- BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding (Devlin et al., 2019) <https://arxiv.org/pdf/1810.04805.pdf>
- <https://towardsdatascience.com/understanding-entity-embeddings-and-its-application-69e37ae1501d>
- <https://www.slideshare.net/AbdallahBashir3/bert-176297542>

BERT fine-tune code:

- <https://skimai.com/fine-tuning-bert-for-sentiment-analysis/>

Comparison of pre-trained models:

- <https://medium.com/@gauravghati/comparison-between-bert-gpt-2-and-elmo-9ad140cd1cda>

Thanks for your attention!